

Unix/Linux System Programming Essentials with Go

Kevser Sırça Alkış

Software Engineer

b!nalyze

DevFest 2024'

Ankara

About Me

- Software Engineer @**Binalyze**
Previously **Hepsiburada, Wope, Digivity, Netinternet**
- Gopher
- Infrastructure & DevOps survivor



 kevsersrca

 kev_src

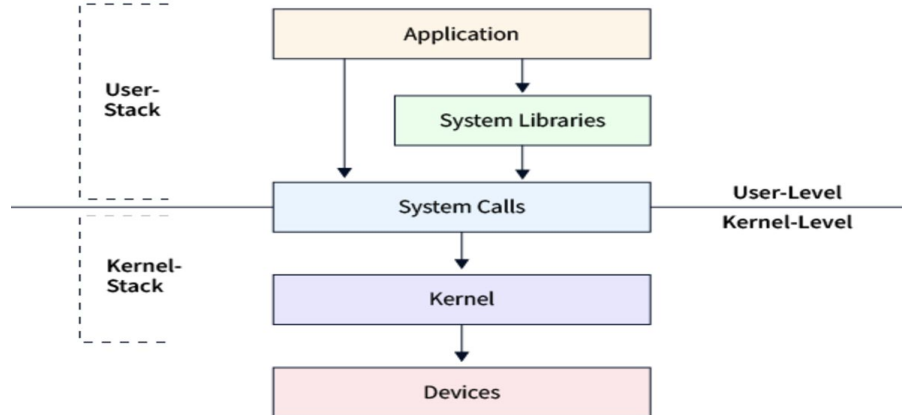
Agenda

1. Definitions
2. Path Manipulations
3. Garbage Collector
4. Processes
5. Exit Codes, Signals and Pipes

Software For Software

System programming (or systems programming) is the process of creating **computer system software**. The definition highlights two main concepts of what system applications are as follows:

- Software that is used by other software, not directly by the final user.



Protection Ring

A **protection ring** is a security architecture used in computer systems to manage access to resources by defining hierarchical privilege levels. It safeguards system resources by preventing unauthorized or faulty access.

Rings Structure:

Ring 0 (Kernel Level): The most privileged level where the operating system kernel operates, with full access to system resources.

Ring 1 and 2: Used for low-level services and device drivers.

Ring 3 (User Level): The least privileged level where user applications run, restricted from direct hardware access.

System Calls

System calls are the way operating systems provide access to the resources for the applications.

It is an **API implemented by the kernel** for accessing the hardware safely.

POSIX Standard

In order to ensure consistency between operating systems, **IEEE formalized some standards** for operating systems.



Windows

Windows is not natively POSIX-compliant. Efforts have been made to bridge the gap.

Open-Source Solutions

- **Cygwin**: Provides a POSIX-like environment on Windows.
- **MinGW**: Supports C applications using the Microsoft Visual C runtime.

Microsoft's Efforts

- **Microsoft POSIX Subsystem**: Early attempt for compliance.
- **Windows Subsystem for Linux (WSL)**:
 - Runs a full Linux environment on Windows.
 - Highly praised for performance and integration.

Linux and macOS

Most Linux distributions follow the **Linux Standard Base (LSB)**, which is another standard that includes POSIX and much more, focusing on maintaining the inter-compatibility between different Linux distributions. It is not considered officially compliant because the developers didn't go into the process of certification.

However, macOS became fully compatible in 2007, and it has been POSIX-certified since then.

Resource Permissions

Mode		Owner	Group	File Size	Last Modified	Filename
drwxrwxrwx	2	sammy	sammy	4096	Nov 10 12:15	everyone_directory
drwxrwx---	2	root	developers	4096	Nov 10 12:15	group_directory
-rw-rw----	1	sammy	sammy	15	Nov 10 17:07	group_modifiable
drwx-----	2	sammy	sammy	4096	Nov 10 12:15	private_directory
-rw-----	1	sammy	sammy	269	Nov 10 16:57	private_file
-rwxr-xr-x	1	sammy	sammy	46357	Nov 10 17:07	public_executable
-rw-rw-rw-	1	sammy	sammy	2697	Nov 10 17:06	public_file
drwxr-xr-x	2	sammy	sammy	4096	Nov 10 16:49	publicly_accessible_directory
-rw-r--r--	1	sammy	sammy	7718	Nov 10 16:58	publicly_readable_file
drwx-----	2	root	root	4096	Nov 10 17:05	root_private_directory

Understanding Mode

drwxrwxrwx

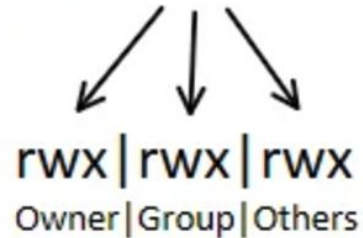
d = Directory

4 r = Read

2 w = Write

1 x = Execute

chmod 777




rwx | rwx | rwx
Owner | Group | Others

7	rwx	111
6	rw-	110
5	r-x	101
4	r--	100
3	-wx	011
2	-w-	010
1	--x	001
0	---	000

640: Owner can read and write and group can read

664: Owner and group can read and write.

Go



fmt

```
func main() {  
    a, b := 3.0, 4.0  
    h := math.Hypot(a, b)  
  
    // Print inserts blanks between arguments when neither is a string.  
    // It does not add a newline to the output, so we add one explicitly.  
    fmt.Print("The vector (", a, b, ") has length ", h, ".\n")  
  
    // Println always inserts spaces between its arguments,  
    // so it cannot be used to produce the same output as Print in this case;  
    // its output has extra spaces.  
    // Also, Println always adds a newline to the output.  
    fmt.Println("The vector (", a, b, ") has length", h, ".")  
  
    // Printf provides complete control but is more complex to use.  
    // It does not add a newline to the output, so we add one explicitly  
    // at the end of the format specifier string.  
    fmt.Printf("The vector (%g %g) has length %g.\n", a, b, h)
```

flag

```
package main
```

```
import (  
    "flag"  
    "fmt"  
)
```

```
func main() {  
    var port int  
    flag.IntVar(&port, "p", 8000, "specify port to use. defaults to 8000.")  
    flag.Parse()  
  
    fmt.Printf("port = %d", port)  
}
```

defer

Deferred function calls are pushed onto a stack. When a function returns, its deferred calls are executed in last-in-first-out order.

```
func main() {  
    fmt.Println("counting")  
  
    for i := 0; i < 10; i++ {  
        defer fmt.Println(i)  
    }  
  
    fmt.Println("done")  
}
```


goroutine

```
package main
```

```
import (  
    "fmt"  
    "time"  
)
```

```
func main() {  
    go func() {  
        for i := 1; i <= 5; i++ {  
            fmt.Println("Hello!")  
            time.Sleep(500 * time.Millisecond)  
        }  
    }()  
  
    for i := 1; i <= 5; i++ {  
        fmt.Println("World!")  
        time.Sleep(700 * time.Millisecond)  
    }  
  
    fmt.Println("Main function finished.")  
}
```

```
World!  
Hello!  
World!  
Hello!  
World!  
Hello!  
World!  
Hello!  
World!  
Main function finished.
```

run, build, install

Command	Purpose	Creates Binary	Installs Binary
<code>go run</code>	Run a program directly	No	No
<code>go build</code>	Compile a program into a binary	Yes	No
<code>go install</code>	Compile and install the binary into <code>\$GOBIN</code>	Yes	Yes

build anywhere

```
GOOS=linux GOARCH=amd64 go build -o app
```

```
GOOS=windows GOARCH=386 go build -o app.exe
```

```
GOOS=darwin GOARCH=arm64 go build -o app
```

build tags

Build tags commonly used to manage platform-specific code, optional features, or dependencies.

```
//go:build linux
```

```
// +build linux      GOOS=linux go build -o app
```

```
package main
```

```
import "fmt"
```

```
func main() {
```

```
    fmt.Println("Running on Linux!")
```

```
}
```

custom build tags

```
//go:build customfeature
```

```
// +build customfeature
```


```
package main
```

```
import "fmt"
```

```
func init() {  
    fmt.Println("Custom feature enabled!")  
}
```

```
go build -tags=customfeature -o app .
```

Path Manipulations



working directory

Linux Command	Description	Go Equivalent	Example Code
<code>pwd</code>	Print the current directory	<code>os.Getwd()</code>	<code>wd, err := os.Getwd()</code>

```
wd, err := os.Getwd()
if err != nil {
    fmt.Println(err)
    return
}
fmt.Println("working dir:", wd)
```

create directory

Linux Command	Description	Go Equivalent	Example Code
<code>mkdir -p <dir></code>	Create nested directories	<code>os.MkdirAll(name string, perm os.FileMode)</code>	<code>err = os.MkdirAll("parent/child", 0o755)</code>
<code>mkdir <dir></code>	Create a directory	<code>os.Mkdir(name string, perm os.FileMode)</code>	<code>err =os.Mkdir("example.txt", 0o755)</code>

rename and move

Linux Command	Description	Go Equivalent	Example Code
<code>mv <src> <dst></code>	Move or rename a file or directory	<code>os.Rename(oldpath, newpath string)</code>	<code>err = os.Rename("old.txt", "new.txt")</code>

```
package main
```

```
import "os"
```

```
func main() {
```

```
    err := os.Rename("kev.txt", "/tmp/output-20241215.txt")
```

```
    if err != nil {
```

```
        panic(err)
```

```
    }
```

```
}
```

path/filepath vs path package

path/filepath

- Handles **local filesystem paths**.
- **Platform-Dependent**: Respects OS-specific separators (\ on Windows, / on Linux).
- Filesystem operations.

path

- Handles **generic or URL paths**.
- **Platform-Independent**: Always uses / as the separator.
- Web paths or cross-platform path manipulation

absolute path and symlinks

```
ls: cannot access 'command': No such file or directory
net2@net2-Laptop:~$ ls -l
total 179860
drwxr-xr-x 2 net2 net2      4096 May 18 06:18 Desktop
drwxr-xr-x 2 net2 net2      4096 May 18 06:18 Documents
drwxr-xr-x 2 net2 net2      4096 May 18 06:18 Downloads
drwxr-xr-x 2 net2 net2      4096 May 18 06:18 Music
lrwxrwxrwx 1 net2 net2           15 Sep  4 04:14 myownlink -> /var/log/syslog
drwxr-xr-x 2 net2 net2      4096 May 18 06:18 Pictures
drwxr-xr-x 2 net2 net2      4096 May 18 06:18 Public
drwx----- 7 net2 net2      4096 Sep  2 07:22 snap
drwxr-xr-x 2 net2 net2      4096 May 18 06:18 Templates
drwxr-xr-x 2 net2 net2      4096 May 18 06:18 Videos
-rw-rw-r-- 1 net2 net2 184133130 Aug 27 23:09 zoom_and64.deb
net2@net2-Laptop:~$
```

absolute path and symlinks

```
func main() {  
    path := "go1.21.13"  
    err := os.Symlink("/usr/local/go/bin/go", path)  
    if err != nil {  
        fmt.Println("Error:", err)  
    }  
  
    if realPath, err := filepath.EvalSymlinks(path); err != nil {  
        path = realPath  
    }  
    fmt.Println("Real Path:", path)  
}
```

absolute path and symlinks

```
package main
```

```
import (  
    "fmt"  
    "path/filepath"  
)
```

```
func main() {  
    if !filepath.IsAbs("test.sh") {  
        if abs, err := filepath.Abs("test.sh"); err == nil {  
            fmt.Println(abs)  
        }  
    }  
}
```

/Users/kevstersirca/go/src/_tmp/test.sh

file create,write and close

```
file, err := os.Create("output.txt")
if err != nil {
    fmt.Println("Error creating file:", err)
    return
}
defer file.Close()

data := "Hello, World!"
n, err := file.Write([]byte(data))
if err != nil {
    fmt.Println("Error writing to file:", err)
    return
}
fmt.Printf("Wrote %d bytes\n", n)
```

Garbage Collector



Memory Management

The operating system handles the **primary** and **secondary** memory usage of the applications. It keeps track of how much of the memory is used, by which process, and what parts are free. It also handles allocation of new memory from the processes and memory de-allocation when the processes are complete.

- Single allocation
- Partitioned allocation
- Paged memory

Go supports Garbage Collection (GC) so you do not have to deal with memory allocation and deallocation.

Virtual Memory

Unix uses the paged memory management technique, abstracting its memory for each application into contiguous virtual memory. It also uses a technique called swapping, which extends the virtual memory to the secondary memory using a swap file

The top command shows details about available memory, swap, and memory consumption for each process:

- RES is the physical primary memory used by the process.
- VIRT is the total memory used by the process, including the swapped memory, so it's equal to or bigger than RES.
- SHR is the part of VIRT that is actually shareable

GC

GC (Garbage Collection) is an automatic memory management mechanism used in programming languages. It identifies memory that is no longer needed (unreachable or unreferenced objects) and reclaims it for reuse.

<https://tip.golang.org/doc/gc-guide>

Marking a Variable for GC

```
package main
```

```
func example() {  
    data := make([]byte, 1024)  
}
```

*// data becomes unreachable after the function ends and
can be collected.*

Marking a Variable for GC

```
package main

func main() {
    var data = make([]byte, 1<<20) // Allocate 1 MB
    ...

    data = nil
    // The data is now eligible for garbage collection
}
```

Trigger the GC Manually

```
package main

import "runtime"

func main() {
    var data = make([]byte, 1<<20) // Allocate 1 MB
    data = nil                       // Remove reference

    runtime.GC()                     // Force garbage collection
}
```

<https://tip.golang.org/doc/gc-guide>

runtime.MemStats

```
func logMemoryUsage() {  
    var memStats runtime.MemStats  
    runtime.ReadMemStats(&memStats)  
    fmt.Printf("Time: %v | Alloc = %v MiB | Sys = %v MiB | NumGC = %v\n",  
        time.Now().Format("15:04:05"),  
        memStats.Alloc/1024/1024, // Current allocation (MiB)  
        memStats.Sys/1024/1024,   // Total allocation (MiB)  
        memStats.NumGC)           // Total count of GC runs  
}
```

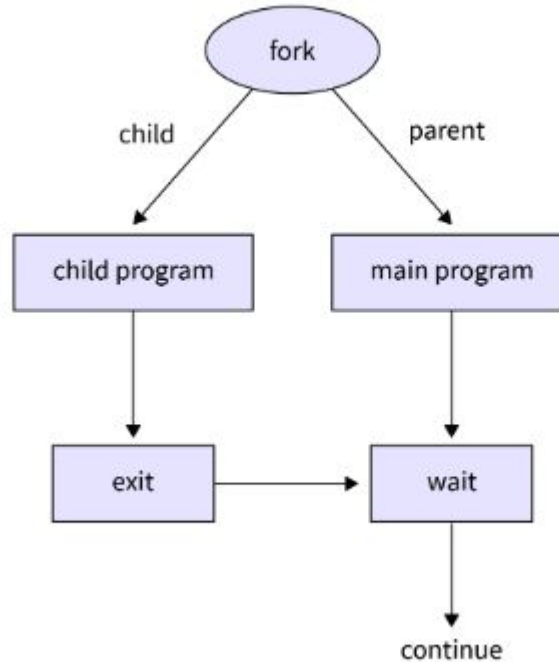
Processes

Processes

When an application is launched, it becomes a process: a special instance provided by the operating system that includes all the resources that are used by the running application.

UID	PID	PPID	C	STIME	TTY	TIME	CMD
root	1	0	0	10:00	?	00:00:01	/sbin/init
user	1234	1	0	10:05	?	00:00:10	go run main.go
user	5678	1234	0	10:10	pts/1	00:00:03	top

Processes Life Cycle



processes

```
func main() {  
    fmt.Println("Current PID:", os.Getpid())  
    fmt.Println("Current Parent PID:", os.Getppid())  
    fmt.Println("User ID:", os.Getuid())  
    fmt.Println("Group ID:", os.Getgid())  
}
```

child process

```
import (  
    "fmt"  
    "os/exec"  
)  
  
func main() {  
    cmd := exec.Command("ls", "-l")  
    if err := cmd.Start(); err != nil {  
        fmt.Println(err)  
        return  
    }  
    fmt.Println("Cmd: ", cmd.Args[0])  
    fmt.Println("Args:", cmd.Args[1:])  
    fmt.Println("PID: ", cmd.Process.Pid)  
    cmd.Wait()  
}
```


stdin and stdout

Standard input can be used to send some data from the application to the child process.
Standard output represents destination sends its output data.

```
b := bytes.NewBuffer(nil)

cmd := exec.Command("cat")
cmd.Stdin = b
cmd.Stdout = os.Stdout
fmt.Fprintf(b, "Hello World! I'm using this memory address: %p", b)
if err := cmd.Start(); err != nil {
    fmt.Println(err)
    return
}
cmd.Wait()
```

Exit Codes, Signals and Pipes



sending exit codes

Applications communicate their result to the operating system by returning a value called **exit status**. This is an integer value passed to the parent process when the process ends.

0: Indicates that the process completed successfully.

Non-zero: Indicates an error or abnormal termination. Different non-zero values often correspond to specific error types.

```
~ > ls non-existing-folder
ls: non-existing-folder: No such file or directory
~ > echo $?
1
~ > █
```

sending exit codes

Exit codes are the way in which a process notifies its parent about its status after terminating.

```
import (  
    "fmt"  
    "os"  
)  
  
func main() {  
    fmt.Println("Hello world!")  
    os.Exit(1)  
}
```

sending exit codes

```
func main() {  
    fmt.Println("Hello, playground")  
    os.Exit(-1)  
}
```

This will have an exit status of 255 even if the function argument is -1 because $(-1) \% 256 = 255$. This happens because the exit code is an 8-bit number (0, 255).

panic

stack unwinding

```
func main() {  
    defer fmt.Println("Hello, playground")  
    panic("panic")  
}
```

reading child process exit code

```
func main() {  
    cmd := exec.Command("ls", "nonexistent-file")  
    err := cmd.Run()  
  
    if err != nil {  
        if exitError, ok := err.(*exec.ExitError); ok {  
            exitCode := exitError.ExitCode()  
            fmt.Printf("Child process exited with code: %d\n", exitCode)  
        }  
    }  
}
```

signals

Exit codes connect processes and their parents, but signals make it possible to interface any process with another, including itself. They are also asynchronous and unidirectional, but they represent communication from the outside of a process.

```
$ kill -l
 1) SIGHUP      2) SIGINT      3) SIGQUIT     4) SIGILL      5) SIGTRAP
 6) SIGABRT     7) SIGEMT     8) SIGFPE     9) SIGKILL     10) SIGBUS
11) SIGSEGV    12) SIGSYS    13) SIGPIPE    14) SIGALRM    15) SIGTERM
16) SIGURG     17) SIGSTOP   18) SIGTSTP   19) SIGCONT    20) SIGCHLD
21) SIGTTIN    22) SIGTTOU   23) SIGIO     24) SIGXCPU    25) SIGXFSZ
26) SIGVTALRM  27) SIGPROF   28) SIGWINCH   29) SIGPWR     30) SIGUSR1
31) SIGUSR2    32) SIGRTMIN  33) SIGRTMIN+1 34) SIGRTMIN+2 35) SIGRTMIN+3
36) SIGRTMIN+4 37) SIGRTMIN+5 38) SIGRTMIN+6 39) SIGRTMIN+7 40) SIGRTMIN+8
41) SIGRTMIN+9 42) SIGRTMIN+10 43) SIGRTMIN+11 44) SIGRTMIN+12 45) SIGRTMIN+13
46) SIGRTMIN+14 47) SIGRTMIN+15 48) SIGRTMIN+16 49) SIGRTMAX-15 50) SIGRTMAX-14
51) SIGRTMAX-13 52) SIGRTMAX-12 53) SIGRTMAX-11 54) SIGRTMAX-10 55) SIGRTMAX-9
56) SIGRTMAX-8  57) SIGRTMAX-7  58) SIGRTMAX-6  59) SIGRTMAX-5  60) SIGRTMAX-4
61) SIGRTMAX-3  62) SIGRTMAX-2 63) SIGRTMAX-1 64) SIGRTMAX
```

signals

```
signalChan := make(chan os.Signal, 1)

signal.Notify(signalChan, os.Interrupt, syscall.SIGTERM)

go func() {
    sig := <-signalChan
    fmt.Printf("Signal catch: %s\n", sig)

    fmt.Println("closing application")
    os.Exit(0)
}()
```

signals

```
p, err := os.FindProcess(pid)
if err != nil {
    panic(err)
}
if err = p.Signal(syscall.SIGTERM); err != nil {
    panic(err)
}
```

```
func os.FindProcess(pid int) (*os.Process, error)
```

FindProcess looks for a running process by its pid.

The Process it returns can be used to obtain information about the underlying operating system process.

On Unix systems, FindProcess always succeeds and returns a Process for the given pid, regardless of whether the process exists. To test whether the process actually exists, see whether `p.Signal(syscall.Signal(0))` reports an error.

`os.FindProcess` on pkg.go.dev


pipes

Pipes are the last unidirectional communication method between processes. As the name suggests, pipes connect two ends – a process input with another process output – making it possible to process on the same host to communicate in order to exchange data.

```
cat input.txt | sort | wc -l
```

pipes

```
reader, writer := io.Pipe()
```



pipes

```
reader, writer := io.Pipe()

go func() {
    defer writer.Close()
    writer.Write([]byte("Hello, Pipe!"))
}()

buf := make([]byte, 20)
n, err := reader.Read(buf)
if err != nil {
    fmt.Println("Error reading from pipe:", err)
    return
}

fmt.Println("Read from pipe:", string(buf[:n]))
```


Thanks!

Kevser Sırça Alkış



@kevsrca



@kev_src